

Week03 - Summary

Data Mining and Hypothesis Testing

- High level overview of statistical tests and data mining (not a deep dive)
- Provide some tools for selecting appropriate statistical tests for evaluating a predictive model, and justifying the choice of tool
- Help you seek details of how to use a statistical method or tool in the data analytic process

Types of Statistical Studies

Observational Studies

1. Retrospective Studies

Example Study: Tanning and Skin Cancer

- The observational study involves 1500 people
- Selected a group of people who had skin cancer and another group who did not have skin cancer
- Asked all participants whether they used tanning beds

2. Prospective Studies

Example Study: Average Computer Time vs Blood Pressure

- Enroll 100 individuals in the observational study
- Ask them to keep track of the computer time they spend each day
- Measure blood pressure

Observational Studies only establish correlation, not causality

Experimental Studies

Example: randomised control trials, A/B tests, drug trials

- Previous “Average Computer Time vs Blood Pressure” as experimental study:
 - 1. Control Group (computer time max 30 minutes)
 - 2. Treatment Group (computer time of at least 2 hours)
 - From the 100 individuals, 50 subjects randomly assigned to each group
 - Factor: average computer time; response: measure blood pressure of each group

Setting up an Environment

- Between subjects:
Each subject sees one and only one condition
- Within subjects:
Subjects see more than one or all conditions

Research Question

- Research question (Q):
 - Asks whether the independent variable has an effect
 - “If there is a change in the independent variable, will there also be a change in the dependent variable?”
- Null hypothesis (H_0):
 - The assumption that there is no effect
 - “There is no change in the dependent variable when the independent variable changes”

Hypothesis Testing

- A hypothesis test examines two opposing hypotheses about a population parameter (e.g. the mean):
 - The null hypothesis
 - The alternative hypothesis
- The null hypothesis represents our initial assumption about the parameter, and we collect evidence to possibly reject the null hypothesis in favour of the alternative hypothesis
 - Example: determine whether the mean of a population differs significantly (this has a special meaning) from a specific value or from the mean of another population

Testing Reliability with p-Values

confidence interval → *Decision*

		Accept H_0	Reject H_0
Truth	H_0 (no difference)	Right Decision	Type I Error
	H_1 (difference exists)	Type II Error	Right Decision

P-value	Indicates	Reject H_0 ?
$\leq \alpha$	Strong evidence against the null hypothesis	Yes
$> \alpha$	Weak evidence against the null hypothesis	No
$= \alpha$	Marginal	NA

Testing which approach is better between subjects

Generate Ratings Data

- We assume different subject groups for each condition

- Each subject sees one of the layouts and is asked to rate on a 5-point Likert scale how strongly he agree or disagree with the statement:
 - Question to subjects: Page gives a good overview?
 - 1=strongly agree; 2=agree; 3=neutral; 4=disagree; 5=strongly disagree
- Outcome:
 - G_data = [1,3,3,2,4,2,3,3,1,5,2,3,4,2,1,3,2,2,1,3,2,3,4,2,1,3,2,2,1,3,1,3,3,2,4,2,3,3,1,5]
 - L_data = [4,5,2,4,4,3,5,4,3,5,1,4,5,3,4,4,2,3,4,5,1,4,5,3,4,4,2,3,4,5,4,5,2,4,4,3,5,4,3,5]
- G_data - grid view
- L_data - list view

Setup: comparing two versions of a display

- Subjects are users of the display (or summary, interface, etc)
 - Dependent variable is user rating (or comprehension, etc)
 - Independent variable is the version of the display
- Q: Do users prefer grid view?
- H_0 : Grid and list data are drawn from the same distribution

Significance: Unpaired Student's t-Test

- Tests the null hypothesis that two population means are equal
- Assumes
 - The samples are independent
 - Populations are normally distributed
 - Standard deviations are equal
- Note:

- Multiply two-tailed p-value by 0.5 for one-tailed p-value
(e.g., to test $A > B$, rather than $A > B$ or $A < B$)

Significance: Mann-Whitney U Test

- Nonparametric version of unpaired t-test
- Assumes:
 - The samples are independent
 - The data is at least ordinal

Testing whether groups differ

Significance: Analysis of Variance (ANOVA)

- Tests the null hypothesis two or more groups have the same population mean
- Assumes:
 - The samples are independent
 - Populations are normally distributed
 - Standard deviations are equal

Significance: Kruskal-Wallis H-test

- Nonparametric version of ANOVA
 - doesn't assume your data comes from a particular distribution such as normal distribution
- Assumes:
 - The samples are independent
- Note:
 - Not recommended for samples smaller than 5
 - Not as statistically powerful as ANOVA

- Both ANOVA and Kruskal-Wallis H-test are extensions of the Mann-Whitney test and Unpaired Student's t-test used to compare the means of more than two populations

Model Evaluation

Measurement: Accuracy, Precision, Recall, F1

	s=1	s=0
g=1	TP (true positives)	FN (false negatives)
g=0	FP (false positives)	TN (true negatives)

g just stands for gold - standard level.

simple specificity / sensitivity table.

- Accuracy: $(TP+TN) / N$
% correct over all instances
- Precision: $TP / (TP+FP)$
% correct system predictions
- Recall: $TP / (TP+FN)$
% correct gold labels
- F1: $2PR / (P+R)$
Harmonic mean of Precision and Recall

Evaluating Classifier Accuracy: Holdout & Cross-Validation Methods

- Holdout method**
 - splits the data randomly into two independent sets
 - Training set (e.g., 2/3) for model construction
 - Test set (e.g., 1/3) for accuracy estimation (also: Validation Set)
 - Random sampling:** a variation of holdout
 - Repeat holdout k times; accuracy = avg. of the accuracies obtained
- Cross-validation** (k-fold, where k = 10 is most popular)

- Randomly partition the data into k mutually exclusive subsets, each approx, equal size
- Leave-One-Out is a particular form of cross-validation:
 - k folds where $k = \#$ of tuples, for small sized data

Data Mining

Types of Statistical Studies: Post-hoc Analysis

Post-hoc Analysis

- testing hypotheses formulated after data collected
- aka data dredging or data mining
- test a wider range of hypotheses faster and cheaper
- may discover surprising patterns
- Be careful with our inferences
 - testing multiple hypotheses
 - control for family-wise error rate
- Confirm findings using experiments
- We'll focus on unsupervised machine learning techniques:
 - Dimensionality reduction
 - **Association rule mining**
 - Clustering
 - Outlier detection
 - etc

Association Rule Mining

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Market-basket transactions

TID: Transaction Identifier

Items: Transactions item set

- Predict occurrence of an item based on other items in the transaction, eg:
 $\{\text{Diaper}\} \rightarrow \{\text{Beer}\}$
 $\{\text{Milk, Bread}\} \rightarrow \{\text{Eggs, Coke}\}$
 $\{\text{Beer, Bread}\} \rightarrow \{\text{Milk}\}$
- Note that arrows indicate co-occurrence, not causality

Definition: Itemset

- An **itemset** is a collection of one or more items
 $\{\text{Milk, Bread, Diaper}\}$
- A **k-itemset** is an itemset containing k items

Definition: Frequent Itemset

- Support count (σ) is the itemset frequency
 $S(\{\text{Milk, Diaper, Beer}\}) = 2$
- Support (s) is the normalised itemset frequency

$$s = \frac{s(\{\text{Milk}, \text{Diaper}, \text{Beer}\})}{|T|} = \frac{2}{5}$$

- A frequent itemset has
 $s \geq \text{min_support}$
- An association rule is an implication of the form $X \rightarrow Y$ where X and Y are itemsets
 $\{\text{Milk}, \text{Diaper}\} \rightarrow \{\text{Beer}\}$
- Confidence (c) measures how often Y occurs in transactions with X

$$c = \frac{s(\{\text{Milk}, \text{Diaper}, \text{Beer}\})}{s(\{\text{Milk}, \text{Diaper}\})} = \frac{2}{3}$$

Mining Association Rules

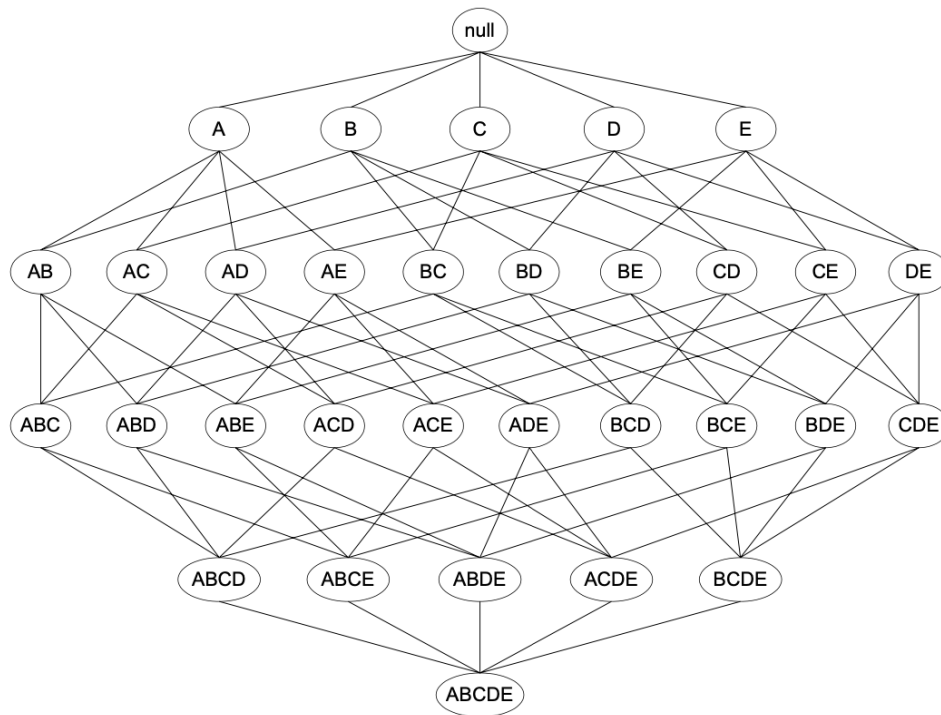
1. Frequent itemset generation

- Generate all itemsets with $s \geq \text{min_support}$

2. Rule generation

- Generate high-confidence rules from each frequent itemset
- Each rule is a binary partitioning of a frequent itemset

There are 2^d candidate itemsets!



Enumeration of 2^5 candidate itemsets for {A, B, C, D, E}

Reducing the number of candidates

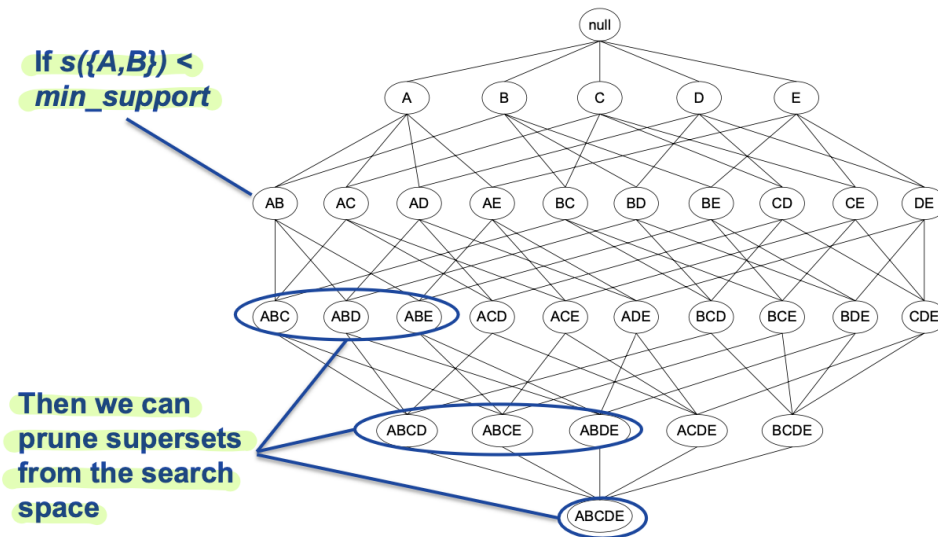
- **Apriori Principle**

If an itemset is frequent, then all of its subsets are also frequent

- Conversely

If an itemset is infrequent, then its supersets are also infrequent

Pruning the 2^d candidate itemsets



Apriori algorithm for generating frequent itemsets

```
While the list of (k-1)-itemsets is non empty:
    Generate candidate k-itemsets
    Identify and keep frequent k-itemsets
```

Create Initial 1-Itemsets

```
def createC1(dataset):
    "Create a list of candidate item sets of size one."
    c1 = []
    for transaction in dataset:
        for item in transaction:
            if not [item] in c1:
                c1.append([item])
    c1.sort()
    # frozenset because it will be a key of a dictionary.
    return list(map(frozenset, c1))
```

Identify itemsets that meet the support threshold

```

def scanD(dataset, candidates, min_support):
    "Returns all candidates that meets a minimum support level"
    sscnt = {}
    for tid in dataset:
        for can in candidates:
            if can.issubset(tid):
                sscnt.setdefault(can, 0)
                sscnt[can] += 1

    num_items = float(len(dataset))
    retlist = []
    support_data = {}
    for key in sscnt:
        support = sscnt[key] / num_items
        if support >= min_support:
            retlist.insert(0, key)
            support_data[key] = support
    return retlist, support_data

```

Generate the next list of candidates

```

def aprioriGen(freq_sets, k):
    "Generate the joint transactions from candidate sets"
    retList = []
    lenLK = len(freq_sets)
    for i in range(lenLK):
        for j in range(i + 1, lenLK):
            L1 = list(freq_sets[i][:k - 2])
            L2 = list(freq_sets[j][:k - 2])
            L1.sort()
            L2.sort()
            if L1 == L2:
                retList.append(freq_sets[i] | freq_sets[j]) # | is set union
    return retList

```

Generate all frequent itemsets

```

def apriori(dataset, min_support = 0.5):
    "Generate a list of candidate item sets"
    C1 = createC1(dataset)
    D = list(map(set, dataset))
    L1, support_data = scanD(D, C1, min_support)
    L = [L1]

```

```

k = 2
while (len(L[k - 2]) > 0):
    Ck = aprioriGen(L[k - 2], k)
    Lk, supK = scanD(D, Ck, min_support)
    support_data.update(supK)
    L.append(Lk)
    k += 1

return L, support_data

```

Identify rules that meet the confidence threshold

```

def calc_confidence(freqSet, H, support_data, rules, min_confidence=0.7):
    "Evaluate the rule generated"
    pruned_H = []
    for conseq in H:
        conf = support_data[freqSet] / support_data[freqSet - conseq]
        if conf >= min_confidence:
            #print(freqSet - conseq, '-->', conseq, 'conf:', conf)
            rules_append((freqSet - conseq, conseq, conf))
            pruned_H.append(conseq)
    return pruned_H

```

Recursively evaluate rules

```

def rules_from_conseq(freqSet, H, support_data, rules, min_confidence=0.7):
    "Generate a set of candidate rules"
    m = len(H[0])
    if (len(freqSet) > (m + 1)):
        Hmp1 = aprioriGen(H, m + 1)
        Hmp1 = calc_confidence(freqSet, Hmp1, support_data, rules,
min_confidence)
        if len(Hmp1) > 1:
            rules_from_conseq(freqSet, Hmp1, support_data, rules, min_confidence)

```

Mine all Association Rules

```

def generateRules(L, support_data, min_confidence=0.7):
    """Create the association rules
    L: list of frequent item sets
    support_data: support data for those itemsets
    min_confidence: minimum confidence threshold
    """
    rules = []
    for i in range(1, len(L)):
        for freqSet in L[i]:
            H1 = [frozenset([item]) for item in freqSet]
            # print("freqSet", freqSet, 'H1', H1)
            if (i > 1):
                rules_from_conseq(freqSet, H1, support_data, rules, min_confidence)
            else:
                calc_confidence(freqSet, H1, support_data, rules, min_confidence)
    return rules

```

Apriori_python

- There are also third party libraries available which implement the apriori algorithm directly
- Example:

```

from apriori_python import apriorir
dataset = [['A', 'C', 'D'], ['B', 'C', 'E'], ['A', 'B', 'C', 'E'],
           ['B', 'E']]
freqItemSet, rules = apriori(dataset, minSup = 0.7, minConf = 0.0)
for r in rules:
    print('{ }==>{ } (c={})'.format(*r))

```